

B-Tree Implementation in C++

Roxanne Farhad

Carrie Hay Kellar

Serena Killion

Design in C++ Final Project

Motivation

- C++ is widely used in database management systems
- B-Trees are one of the most important data structures for DBMS
- No standard library container representing B-Trees



Why is there no STL tree container already?

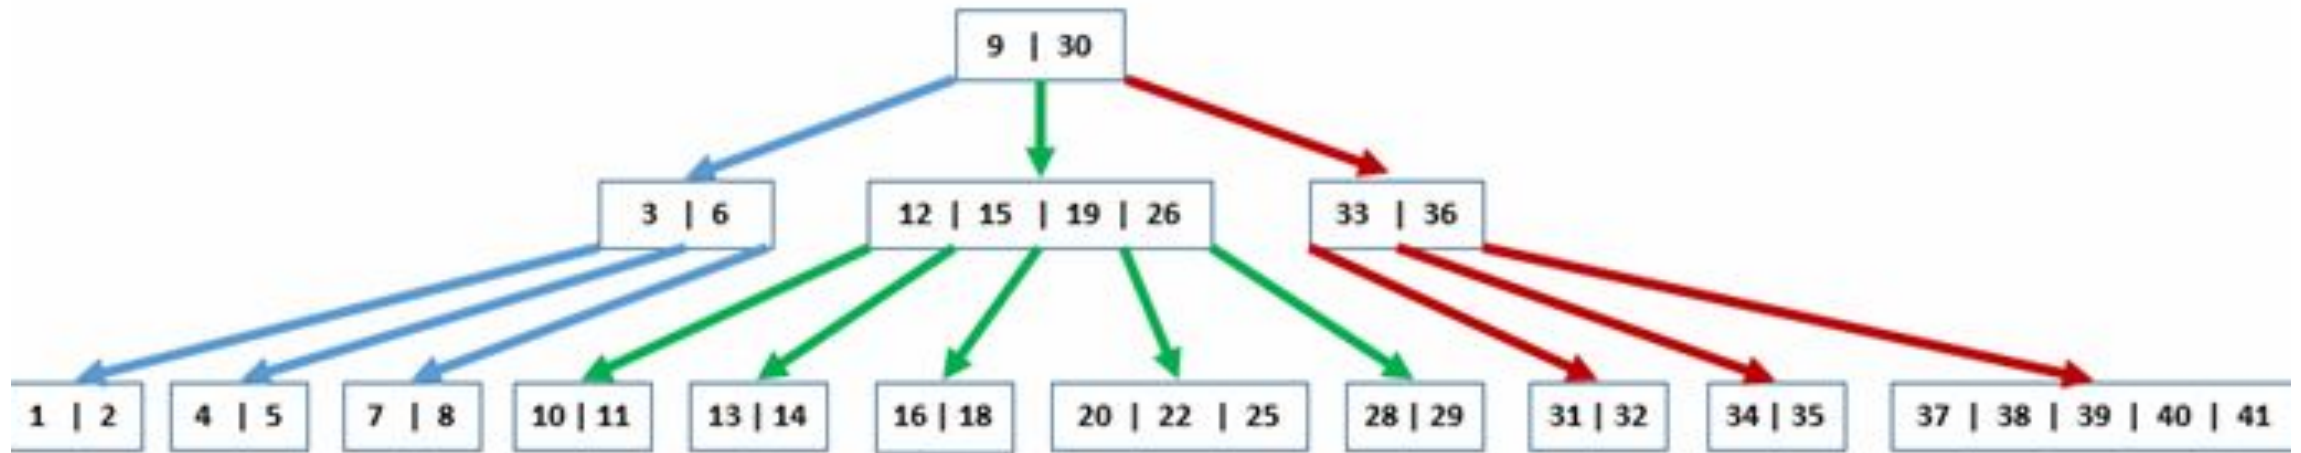
- Containers defined by interface and observable characteristics not underlying structure
 - `std::map` and `std::set` are internally implemented as red-black trees but tree functionality is not exposed
- Difficult to have a general tree container that satisfies everyone
 - How to decide on iterator?
 - Is the number of children per node fixed?
 - How much pointer overhead per node?
 - Which algorithms to include?

Why B-Tree?

- specialized m-way search tree which is a BST with multiple elements in a node
- sorted data structure is optimal for searching, inserting, and deleting
- get all keys using inorder traversal
- natural to find closest lower or greater elements or range queries
- all operations guaranteed to work in log time
 - hash tables has constant average time, but table resizing is costly

Type	Insertion	Deletion	Lookup
Unsorted Array	$O(1)$	$O(n)$	$O(n)$
Sorted Array	$O(n)$	$O(n)$	$O(\log(n))$
B-tree	$O(\log(n))$	$O(\log(n))$	$O(\log(n))$


Properties of B-Tree



- Ordered, keys within node are ordered too
- Non leaf nodes must have $n-1$ keys and n child pointers
- All leaf nodes are at same level
- All non-root nodes have between $\lceil m/2 \rceil - 1$ keys and $2m-1$ keys
- Non-leaf root node must have at least 2 children

Application to Database Indexing

- B-Trees are optimized for data stored on disk
- Average read operator of a sector: 15 ms
- Node size usually correlates to disk block size
- Reduces number of disk risks → faster access



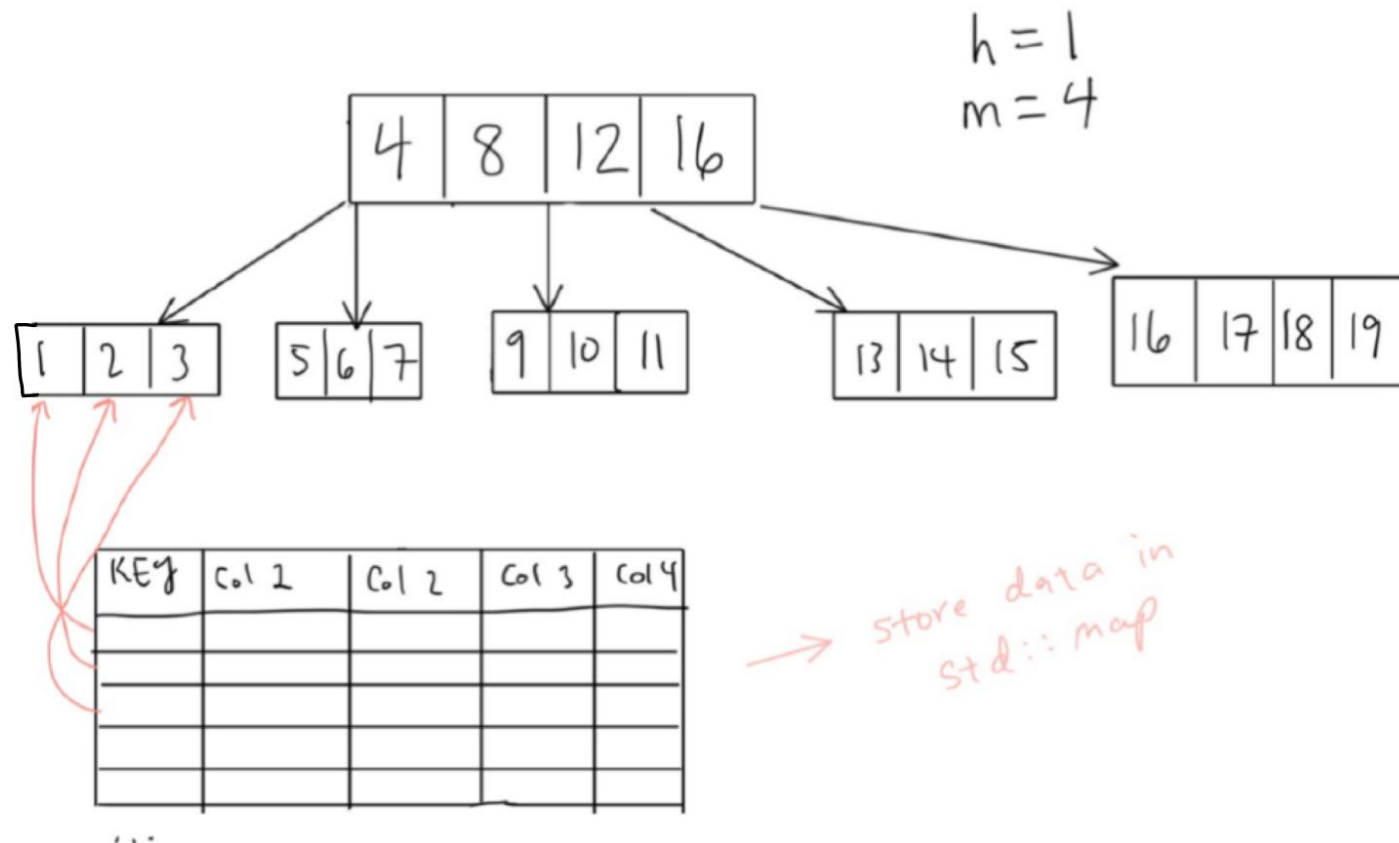
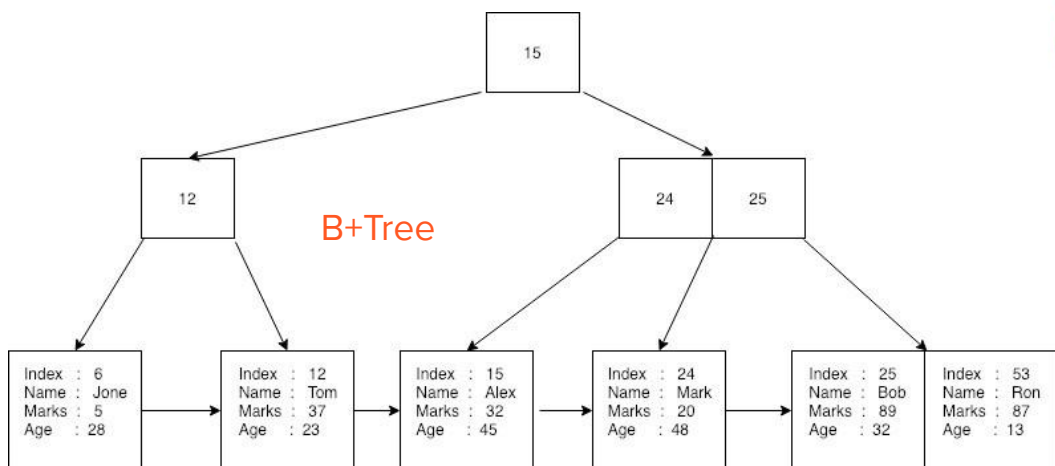
From a practical point of view, B-trees, therefore, guarantee an access time of less than 10 ms even for extremely large datasets.

—Dr. Rudolf Bayer, inventor of the B-tree


Our Implementation

Many DBMS use b+tree, a b-tree with an additional level of linked leaves with copies of key and data records.

We wanted to focus on the b-tree so for the real-world data we stored the key (date) in the nodes and the key-values in a `std::map`.



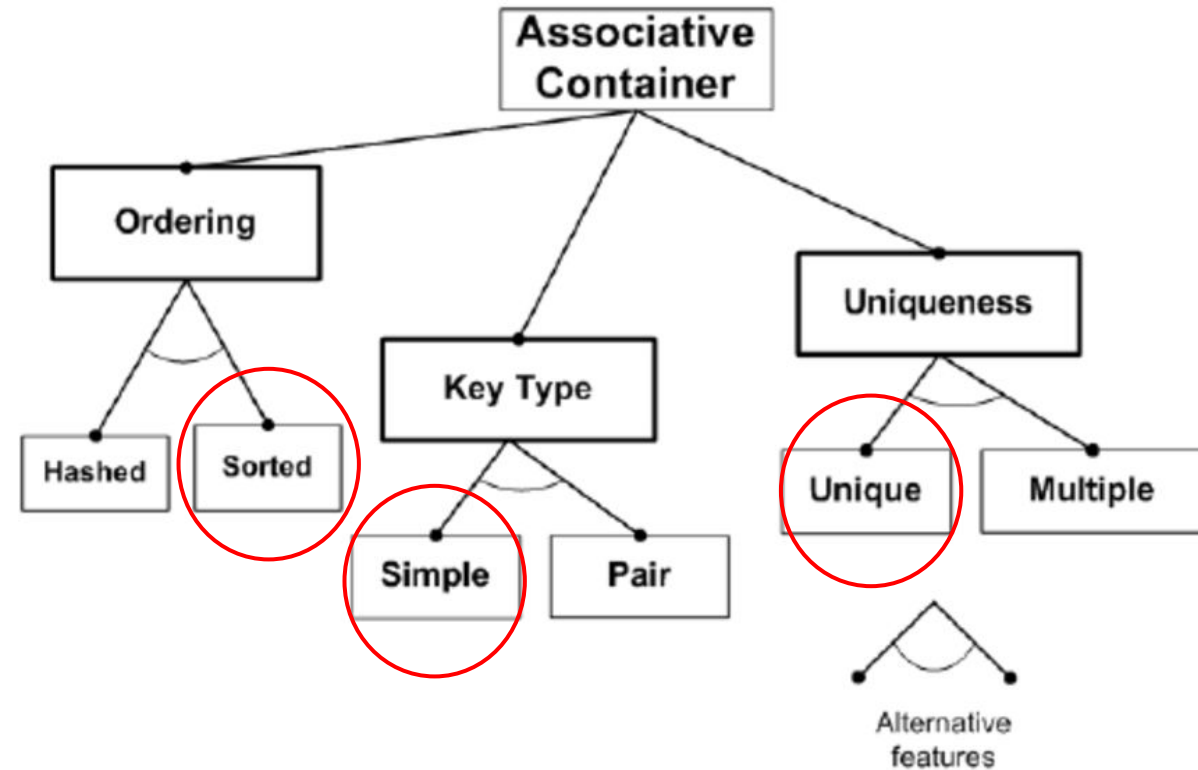
Demo



Designing BTree

- Create a btree based on the STL container standard for associative containers
- 3 main parts to a container:
 - Following STL shared properties (focus)
 - Allocators (V2)
 - Iterators (V2)
- Implemented in class templates

Feature diagram for Associative Container



BTree & BNode class

```
template<typename T>
class BTree {
    using btree_type = BTree<T>;
private:
    BNode<T> *root;

    int m; /* pre-defined degree of tree */
    int (*compare)(T, T); /* comparative function for tree (as templated) */
    void (*printKey)(T); /* Function used to print items in the tree. */
    int index;
    int elemCount;

public:
    BTree( int, int (*)(T, T), void (*)(T) ); /* constructor with defined type */
    ~BTree( ); /* destructor */

    void clear(); /* clears the tree */
    int insert(T); /* inserts a key into the tree */
    T remove(T); /* removes a key from tree */
    int size(); /* returns the size of the tree */
    bool empty(); /* returns bool if the tree is empty or not */

    BNodeKey<T>* search(T); /* returns the node associated with key */
    void traverse(); /* traverses the tree */
};
```

```
template<typename T>
class BNode
{
    using BNodeType = BNode<T>;
public:
    bool isLeaf;
    int m; /* order of the tree */
    int size; /* refers to the number of keys */
    std::vector< BNodeKey<T> > *keys;
    std::vector<BNodeType > *children;

private:
    int (*compare)(T, T);
    void (*printKey)(T);

public:
    BNode(int m, int (*)(T, T), void (*)(T)); /* constructor */
    ~BNode( ); /* destructor */
    void print( );
    unsigned insertKey(T k, int index); /* inserts the key into node */
    unsigned insertChild(int i, BNodeType *x);
    void removeKey(int i, int j);
    void removeChild(int i, int j);
};
```

User Interface Design

COMMANDS:

1. `SELECT <PK>, ..., <columnN> FROM table [WHERE <PK> = x]`
2. `INSERT (<PK>, ..., <columnN>) (<PK>, ..., <valueN>)`
3. `DELETE WHERE PK = x`

1. **Parse** a data file from the user to create a table
 - a. The PK of each row is inserted into the btree and a unique index corresponding to that PK is returned
 - b. The index of the PK is then inserted into a map and corresponds to a vector of string values that represent the associated data for that row.
2. **Select** grabs the index from the btree and then finds the resulting vector from the map using that index.
3. **Delete** returns the index of the key in the node while deleting it from the btree, and then uses that index to remove the values from the map.

```
SELECT * FROM data
```

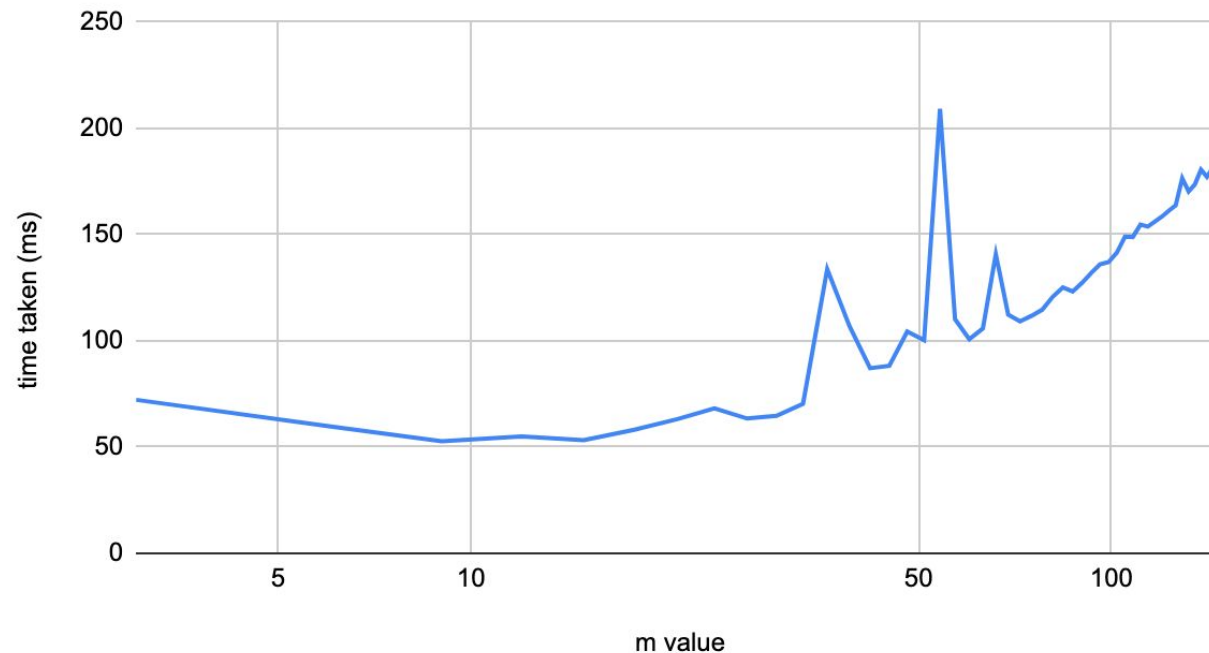
artist	date	last-week	peak-rank	rank	song	weeks-on-board
Ricky Nelson	04/08/1958		1	1	Poor Little Fool	1
Ricky Nelson	11/08/1958	1	1	1	Poor Little Fool	2

Tests

- As we are designing an associative STL container, we wanted to compare our functions to current maps and sets that use red-black trees for their operations.
- Tested insert by timing how long it would take each containers to insert a range from 50,000 to 1 million random values.
- Tested search by choosing 10 random values within a range of 50,000 to 1 million values, and averaged the time it took each container to find the value.

Measurement Results

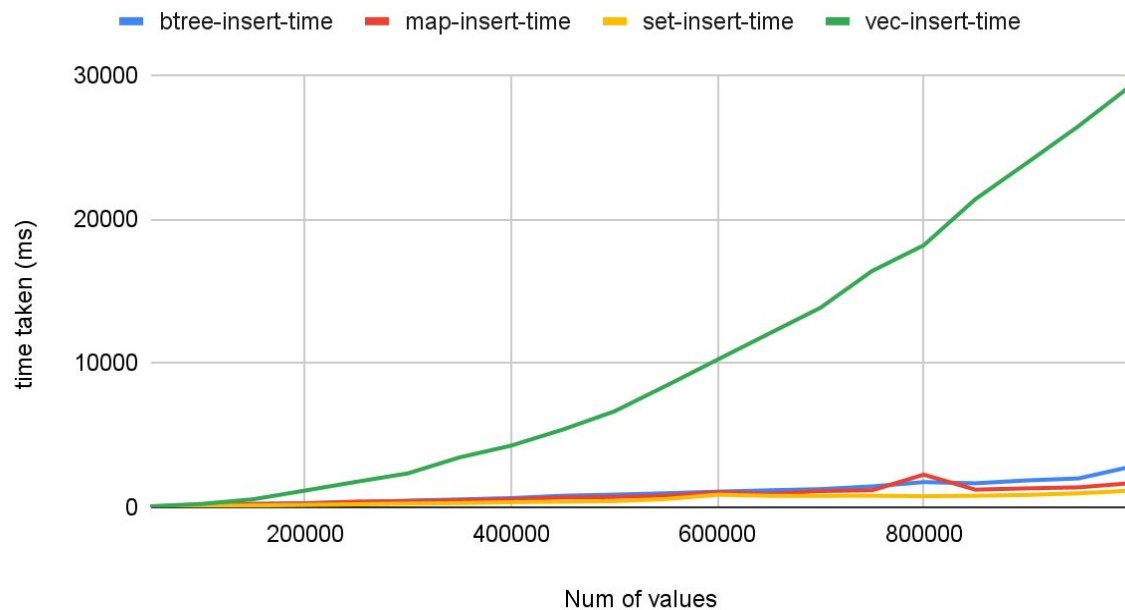
- m: minimum degree
- insertion of 50,000 random telements
- larger value of m → wider, shorter tree




Measurement Results

- B-Tree inserts and searches faster than Vector

insert(), m = 12




search(), m = 12



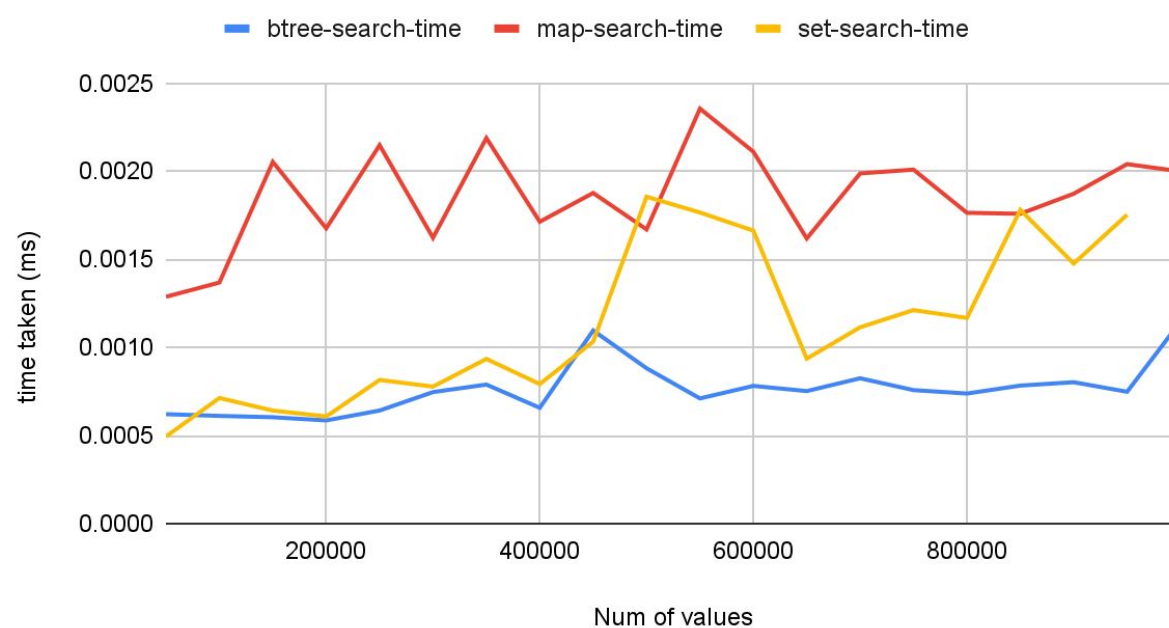
Measurement Results

- B-Tree faster at searching, roughly similar for inserting

insert(), m = 12



search(), m = 12



Conclusion

“If I were designing STL today, I would have a different set of containers. For example, an in-memory B^ -tree is a far better choice than a red-black tree for implementing an associative container.”*

- Alexander Stepanov

V.20

- Implement B+ Tree
- SQL interface parser
- Multi-lookups
- Iterator and allocator implementations